

Exhibit*Guide*

A smart audio & content guide for museums and galleries. Visitors scan the exhibit label with their own phone — no hardware rental, no app install, no account — and get the story behind the work, their way.

Progressive Web App

QR + NFC triggered

B2B2C — institutions subscribe

I. SCAN

QR code or NFC tag on any exhibit label opens the page instantly.

II. READ OR LISTEN

A 30-second TLDR up front; text, audio, video and imagery on demand.


III. CATALOGUE

Optionally save exhibits to a personal collection to revisit after the visit.

JOHANNES VERMEER

*Girl with a Pearl
Earring, c. 1665*

Oil on canvas · 44.5 × 39 cm



Scan for the story —

*thirty seconds, or as
deep as you like.*

Standing in front of a masterpiece shouldn't require *a hardware rental*.

A VISITOR NEEDS TO...

- a.* **Understand what they're looking at** — context, history, technique — without queueing for a handset.
- b.* **Control the depth** — thirty seconds for a tourist, full provenance for an enthusiast.
- c.* **Keep what they found** — a record of the visit, ready to cite after leaving the building.

WHAT THEY USE TODAY — AND THE GAP

Rented audio handsets	Costly hardware, queues, hygiene concerns, fixed linear tours.
Native museum apps	Nobody pauses mid-visit to download an app and create an account.
Printed wall labels	Universal but shallow — no audio, no depth, nothing to take home.

THE USP

No current system combines frictionless access, genuine content depth, and a personal catalogue. ExhibitGuide fills that hole in the market.

Museum & gallery visitors **PRIMARY — THE USER**

Casual tourists through academic researchers — phone already in hand, on patchy gallery WiFi.

Museum & gallery operators **SECONDARY — THE PAYING CUSTOMER**

Curators publishing rich content without engineering support — a flat-fee SaaS subscription per institution.

Prototype fast, release to *real visitors*, and let the numbers judge.

i. Rapid prototyping

Short cycles: wireframes → clickable prototype → working increment, with user feedback gathered at the end of every weekly sprint.

ii. User Acceptance Testing

Release the MVP to real users in a real space — a pilot with one small independent gallery, where patchy WiFi, glare and mixed devices surface the true friction points.

iii. Qualitative validation

Visitor interviews and short exit surveys alongside the analytics — does the scan moment feel effortless, and is the content worth it?

CORE SUCCESS CRITERION

Scan → content
in *< 2 seconds*

Live gallery test, iPhone
Safari + Android Chrome

ADOPTION

User sign-ups *per week*

Tracked via auth / database
records

ENGAGEMENT

Daily active *users*

Tracked via app analytics

RETENTION

7-day retention
rate

Tracked via app analytics

Agile in one-week sprints, on a stack chosen for *friction-free delivery*.

Development method **AGILE — SCRUM FLAVOUR**

Wireframes mapped to user stories; a prioritised feature backlog; a demoable increment presented every Thursday.

Version control **GITHUB**

Feature-branch workflow — each branch maps to a backlog item and sprint deliverable; pull-request review; main always deployable.

DEVELOPMENT SCHEDULE — FOUR ITERATIVE CYCLES

TUE Kick-off Initial project understanding session	SPRINT 01 Thu–Thu Present 10 min + artefacts: wireframes & backlog	SPRINT 02 Thu–Thu + first app parts: scan → exhibit content page	SPRINT 03 Thu–Thu + current app: CMS, accounts, format preference	SPRINT 04 Thu–Thu + near-complete app, UAT findings & polish
--	--	--	---	--

TECHNOLOGY STACK — WHY EACH LAYER

<i>Frontend</i> React PWA + Tailwind No app install; offline-capable; mobile-first by default.	<i>Triggers</i> Camera QR + Web NFC QR works on every phone; NFC is an Android-only bonus.	<i>Backend</i> Django + DRF Python suits LLM integration; clean API for app and CMS.	<i>Data</i> PostgreSQL + Redis Content, accounts, collections; cached for fast page loads.	<i>Media</i> S3 storage + CDN Audio, video and images served fast on patchy gallery WiFi.	<i>Intelligence</i> Claude API + Celery Phase 3 TLDRs & translation, generated async, cached once.
---	---	---	---	--	---

Three phases — the MVP is *this module's build target*.

PHASE 1 MVP — BUILD TARGET

Scan to story

Frictionless by default — no account needed

- **QR scanning** — phone camera opens the PWA page directly
- **NFC tap** — same page via tag, on Android Chrome
- **Exhibit pages** — TLDR + expandable full content
- **Format preference** — text, audio narration, or both
- **Rich media** — images, audio player, video embeds
- **Museum CMS** — curators publish without engineers
- **Optional accounts** — login / logout, only for cataloguing

PHASE 2

Personal collections

For researchers and returning visitors

- **One-tap save** — catalogue any exhibit while walking
- **Collection view** — full content accessible post-visit
- **Notes** — attach your own thoughts to saved pieces
- **Export** — PDF or plain text, citation-ready
- **Visit history** — passive log of scanned exhibits

PHASE 3

Intelligence layer

AI assistance for curators and visitors

- **AI summaries** — Claude drafts TLDRs from curator data
- **Related exhibits** — thematic connections surfaced
- **Multi-language** — content translated on demand
- **Accessibility mode** — simplified language, richer audio
- **Operator analytics** — engagement & dwell per exhibit

Six people, one gallery, *one scan each*.

QR DEEP-LINK · NO ACCOUNT

As a first-time visitor,

I WANT exhibit pages to open straight from my camera,

SO THAT I never install an app or sign up mid-visit.

TLDR SUMMARY

As a time-pressed tourist,

I WANT a thirty-second summary the moment I scan,

SO THAT I get the story without breaking my stride.

AUDIO · FORMAT PREFERENCE

As a visitor absorbed in the work,

I WANT narration to play in my earphones,

SO THAT I learn without taking my eyes off the piece.

COLLECTIONS · NOTES · EXPORT

As an academic researcher,

I WANT to save exhibits with my own notes and export them,

SO THAT I can cite what I saw in my writing.

LOGIN · VISIT HISTORY

As a returning member,

I WANT to log in and find every past visit,

SO THAT I pick up exactly where I left off.

MUSEUM CMS

As a curator,

I WANT to publish exhibit content from a simple CMS,

SO THAT I update the gallery without engineering help.

Measure the *whole funnel*, with tooling that respects visitor privacy.

THE KPIS THAT MATTER — ACROSS THE VISITOR FUNNEL

ACQUISITION Sign-ups <i>per week</i> Target: growth WoW Source — auth records	ACTIVATION Scan → content Target: < 2 s Source — Sentry / Web Vitals	ENGAGEMENT Daily active users · <i>format mix</i> Target: rising DAU Source — PostHog	RETENTION 7-day <i>retention rate</i> Target: ≥ 30% Source — PostHog cohorts	OPERATOR VALUE Scans & dwell <i>per exhibit</i> Target: institution KPI Source — operator dashboard
---	--	---	--	---

THE ANALYTICS STACK — AND WHY EACH

<i>Product analytics</i> PostHog	Funnels, retention cohorts & session replay. Self-hostable, so visitor data stays in-house per tenant.	<i>Traffic</i> Plausible	Cookieless and GDPR-clean — no consent banner to interrupt the scan-to-content moment.	<i>Performance</i> Sentry + Web Vitals	Tracks scan-to-content latency and errors on real devices, guarding the < 2 s bar.
<i>Event pipeline</i> Django → PostgreSQL	Anonymised scan, dwell and format events — the source of truth behind operator analytics.	<i>Operator BI</i> Metabase	Self-hosted dashboards on the event store; curators read per-exhibit insight without SQL.	<i>CI budgets</i> Lighthouse	Performance budgets in the pipeline keep the PWA under the 2-second target before every release.